



# TENSORFLOW

...

Introduction Session

# Choice of Toolkit

- Model specification:
  - Configuration file (e.g. Caffe, DistBelief, Microsoft CNTK) versus programmatic generation (e.g. Torch, Theano, Tensorflow)
- Choice of high-level language:
  - Lua (Torch) vs. Python (Theano, Tensorflow) vs C++(Caffe/Tensorflow)
- People work with python because of rich community and library infrastructure
  - Shift of Torch to PyTorch, caffe libraries in python

# Why Tensorflow

- Visualization (TensorBoard is da bomb)
- Checkpoints (for managing experiments)
- Developed and maintained by Google Brain
- Auto-differentiation autodiff (no more taking derivatives by hand. Yay)
- Used by Google, OpenAI, DeepMind, Snapchat, Uber, Airbus, eBay, Dropbox, etc

# How Tensorflow

- “Tensors”: multi dimensional arrays
  - They flow thru a “computational graph”. Duh, that’s why the name
- 2 stages of writing TF code
  - Building the graph
  - Running the graph

# How Tensorflow

- “Tensors”: multi dimensional arrays
  - They flow thru a “computational graph”. Duh, that’s why the name
- 2 stages of writing TF code
  - Building the graph
  - Running the graph
- Most Important Point when trying first time: Above Point

# How Tensorflow

- “Tensors”: multi dimensional arrays
  - They flow thru a “computational graph”. Duh, that’s why the name
- 2 stages of writing TF code
  - Building the graph
  - Running the graph
- Most Important Point when trying first time: Above Point
- Most Important Point every other time: Two points above

# What is TensorFlow ?

TensorFlow is a way of representing computation without actually performing it until asked

- Avoids doing a lot of unnecessary computation
- Allows good parallelization of computation
- Handles on-line data (non stop) pretty well
- Executes only relevant sub-graphs (auto-generated) based on need



# Numpy vs TensorFlow

(Source: Stanford Course)



# Simple Numpy Recap

```
In [23]: import numpy as np
```

```
In [24]: a = np.zeros((2,2)); b = np.ones((2,2))
```

```
In [25]: np.sum(b, axis=1)
```

```
Out[25]: array([ 2.,  2.])
```

```
In [26]: a.shape
```

```
Out[26]: (2, 2)
```

```
In [27]: np.reshape(a, (1,4))
```

```
Out[27]: array([[ 0.,  0.,  0.,  0.]])
```

# Repeat in TensorFlow

More on *Session*  
soon

```
In [31]: import tensorflow as tf
```

```
In [32]: tf.InteractiveSession()
```

```
In [33]: a = tf.zeros((2,2)); b = tf.ones((2,2))
```

```
In [34]: tf.reduce_sum(b, reduction_indices=1).eval()
```

```
Out[34]: array([ 2.,  2.], dtype=float32)
```

More on *.eval()*  
in a few slides

```
In [35]: a.get_shape()
```

```
Out[35]: TensorShape([Dimension(2), Dimension(2)])
```

*TensorShape* behaves  
like a python tuple.

```
In [36]: tf.reshape(a, (1, 4)).eval()
```

```
Out[36]: array([[ 0.,  0.,  0.,  0.]], dtype=float32)
```

# Common Numpy Functions

```
np.zeros()  
np.ones()  
np.linspace()  
np.range()  
np.mul()  
np.mean()  
np.sum()
```

---

# Common Numpy Functions

And their TF equivalents

<code>np.zeros()</code>	<code>⇒tf.zeros()</code>
<code>np.ones()</code>	<code>⇒tf.ones()</code>
<code>np.linspace()</code>	<code>⇒tf.linspace()</code>
<code>np.range()</code>	<code>⇒tf.range()</code>
<code>np.mul()</code>	<code>⇒tf.mul()</code>
<code>np.mean()</code>	<code>⇒tf.reduce_mean()</code>
<code>np.sum()</code>	<code>⇒tf.reduce_sum()</code>

---

# Common Numpy Functions

And their TF equivalents

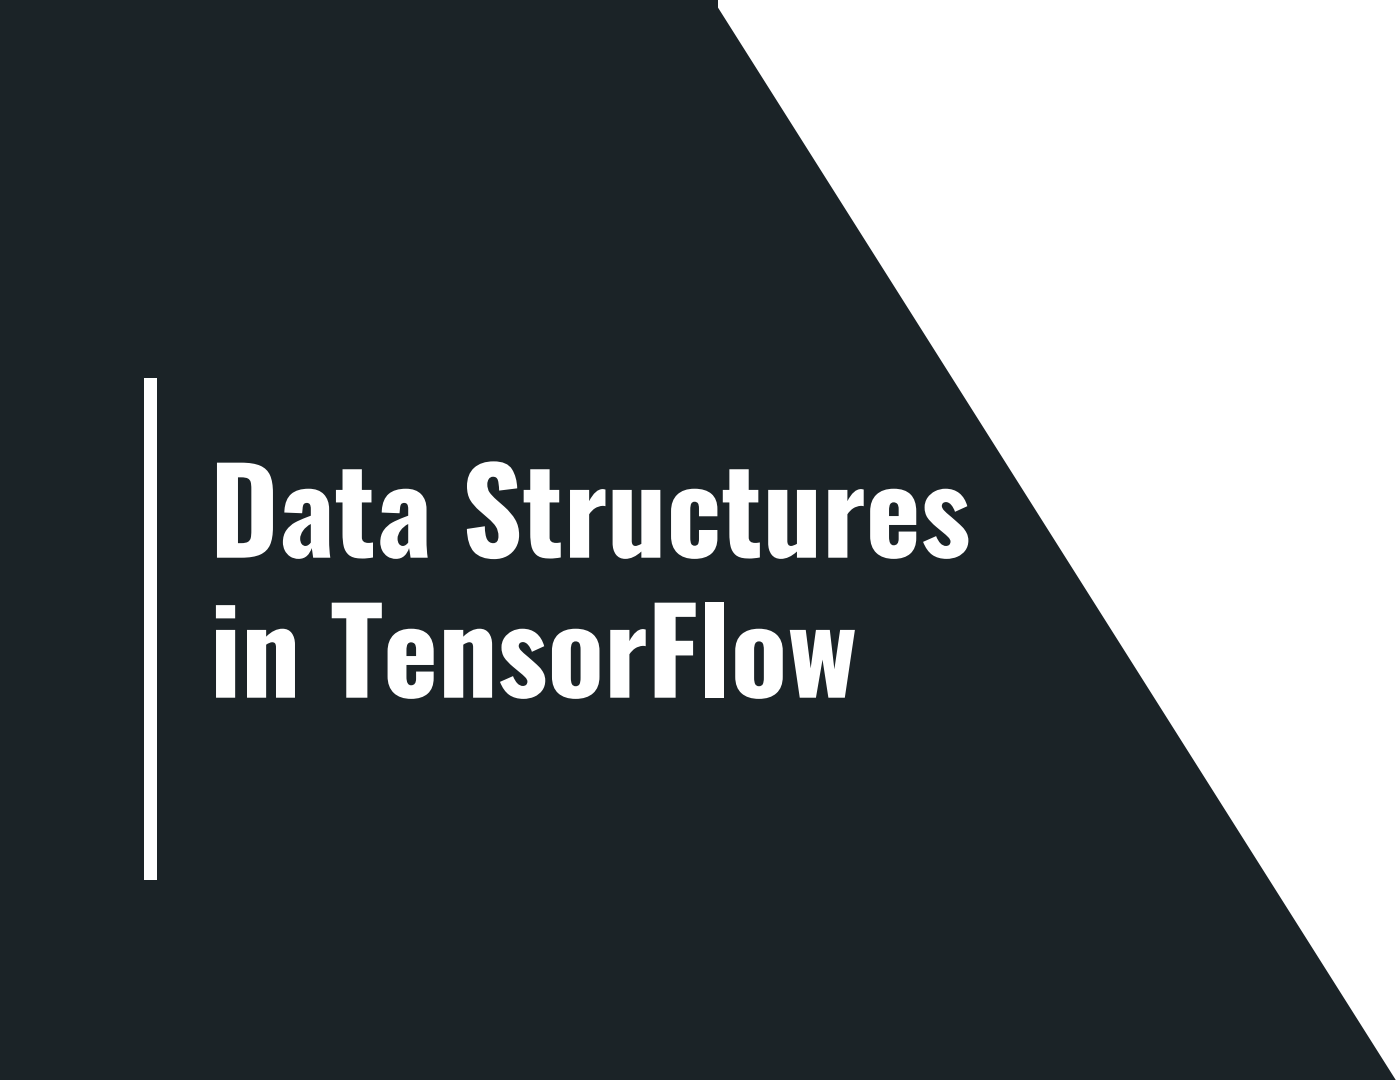
More or less have same syntax

<code>np.zeros()</code>	<code>⇒tf.zeros()</code>
<code>np.ones()</code>	<code>⇒tf.ones()</code>
<code>np.linspace()</code>	<code>⇒tf.linspace()</code>
<code>np.range()</code>	<code>⇒tf.range()</code>
<code>np.mul()</code>	<code>⇒tf.mul()</code>
<code>np.mean()</code>	<code>⇒tf.reduce_mean()</code>
<code>np.sum()</code>	<code>⇒tf.reduce_sum()</code>

---

# TF vs np

- TensorFlow integrates seamlessly with NumPy
  - `tf.int32 == np.int32 # True`
- Can pass numpy types to TensorFlow ops
  - `tf.ones([2, 2], np.float32) # ⇒ [[1.0 1.0], [1.0 1.0]]`
  - Exception: `np.float128` doesn't exist in tensorflow, as people normally use lower precision floats for speed up
- For `tf.Session.run(fetches)`: If the requested fetch is a Tensor, then the output of will be a NumPy ndarray.



# **Data Structures in TensorFlow**

# Data structures in TensorFlow

- `tf.constant`: Inputs as pre-defined constants
- `tf.Placeholder`: Inputs of variable size
- `tf.Variable`: Variables of fixed size



# Data structures in TensorFlow

- `tf.constant`: Inputs as pre-defined constants
- `tf.Placeholder`: Inputs of variable size
- `tf.Variable`: Variables of fixed size

Often, you will feed in the dataset finally when executing the graph. The dataset can contain inputs of a previously undefined size. A placeholder is a promise to provide a value as an input later.

# Data structures in TensorFlow

- `tf.constant`: Inputs as pre-defined constants
- `tf.Placeholder`: Inputs of variable size
- `tf.Variable`: Variables of fixed size

To make the model trainable, we need to be able to modify the graph to get new outputs with the same input. Variables allow us to add trainable parameters to a graph.

Ex: Weights and Biases of NN

# Data structures in TensorFlow

- `tf.constant`: Inputs as pre-defined constants
- `tf.Placeholder`: Inputs of variable size
- `tf.Variable`: Variables of fixed size

Constants: initialized in definition (`a = tf.constant()`)

Variables: Not initialized when you call `tf.Variable`.

To initialize all the variables, run

```
>> init = tf.global_variables_initializer()  
>> sess.run(init)
```

# Example Code: (run it)

```
import tensorflow as tf
```

```
x1 = tf.constant(35)
```

```
x2 = tf.constant(5)
```

```
y = tf.add(x1,x2)
```

```
print(y)
```

---

# Example Code: (run it)

How do we get the output then?

```
import tensorflow as tf
```

```
x1 = tf.constant(35)
```

```
x2 = tf.constant(5)
```

```
y = tf.add(x1,x2)
```

```
print(y)
```

---

# How do we get the output then?

We only built the graph.  
We didn't execute it

```
import tensorflow as tf
```

```
x1 = tf.constant(35)
```

```
x2 = tf.constant(5)
```

```
y = tf.add(x1,x2)
```

```
model =
```

```
tf.global_variables_initializer()
```

```
with tf.Session() as sess:
```

```
    sess.run(model)
```

```
    print(sess.run(y))
```

# Running a graph

- Two general methods:
  - `sess.run()`
  - `<variable_name>.eval()`
- Both do the same thing. Reason for using `run()` over `eval`:  
You can evaluate multiple variables at the same time  
(Do `sess.run([a,b,c,..])`)

# tf.Variable.assign()

```
W = tf.Variable(10)
W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    print W.eval() # >> ????
```

What will the output be?



# tf.Variable.assign()

```
W = tf.Variable(10)
W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    print W.eval() # >> 10
```

Woah. What the ...?

# tf.Variable.assign()

```
W = tf.Variable(10)
W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    print W.eval() # >> 10
```

W.assign(100) doesn't assign the value 100 to W. It creates an assign op, and that op needs to be run to take effect.

# tf.Variable.assign()

```
W = tf.Variable(10)
W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    print W.eval() # >> 10
```

-----

```
W = tf.Variable(10)
assign_op = W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    sess.run(assign_op)
    print W.eval() # >> 100
```

# tf.Variable.assign()

```
W = tf.Variable(10)
W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    print W.eval() # >> 10
```

You don't need to initialize variable  
because assign\_op does it for you

```
-----
W = tf.Variable(10)
assign_op = W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    sess.run(assign_op)
    print W.eval() # >> 100
```

```
W = tf.Variable(10)
assign_op = W.assign(100)
with tf.Session() as sess:
    sess.run(assign_op)
    print W.eval() # >> 100
```

# Visualising Graph

```
import tensorflow as tf

x = tf.constant(35, name='x')
y = tf.Variable(x + 5, name='y')
model = tf.global_variables_initializer()

with tf.Session() as sess:
    sess.run(model)
    print(sess.run(y))
    writer = tf.summary.FileWriter('/tmp/tfGraphs', sess.graph)

writer.close() # close the writer when you're done using it
```

# Visualising Graph

```
import tensorflow as tf
```

```
x = tf.constant(35, name='x')
```

```
y = tf.Variable(x + 5, name='y')
```

```
model = tf.global_variables_initializer()
```

```
with tf.Session() as sess:
```

```
    sess.run(model)
```

```
    print(sess.run(y))
```

```
    writer = tf.summary.FileWriter('/tmp/tfGraphs', sess.graph)
```

```
writer.close() # close the writer when you're done using it
```

Run with  
tensorboard --logdir=/tmp/tfGraphs



Fit to screen



Download PNG

Run

(1)

Session

runs (0)

Upload

Choose File

Trace inputs



Color



Structure



Device

colors

same substructure



unique substructure

Graph

(\* = expandable)




Namespace\*



OpNode

## Main Graph



## Auxiliary Nodes



# Data Structures

## Placeholder

- Commonly used for input to graph
- Requires feeding of value for the variable at graph run time, rather than variable declaration time.
- Data fed in using `feed_dict` in graph run command



# Placeholders

Note the usage of 'None' to specify undefined size for the input

```
import tensorflow as tf

x = tf.placeholder("float", [None, 3])
y = x * 2

with tf.Session() as session:
    x_data = [[1, 2, 3],
               [4, 5, 6],]
    result = session.run(y,
                           feed_dict={x: x_data})
    print(result)
```


---




# Getting Stuff Done In TF

(Taken from NLintz' Tutorials)

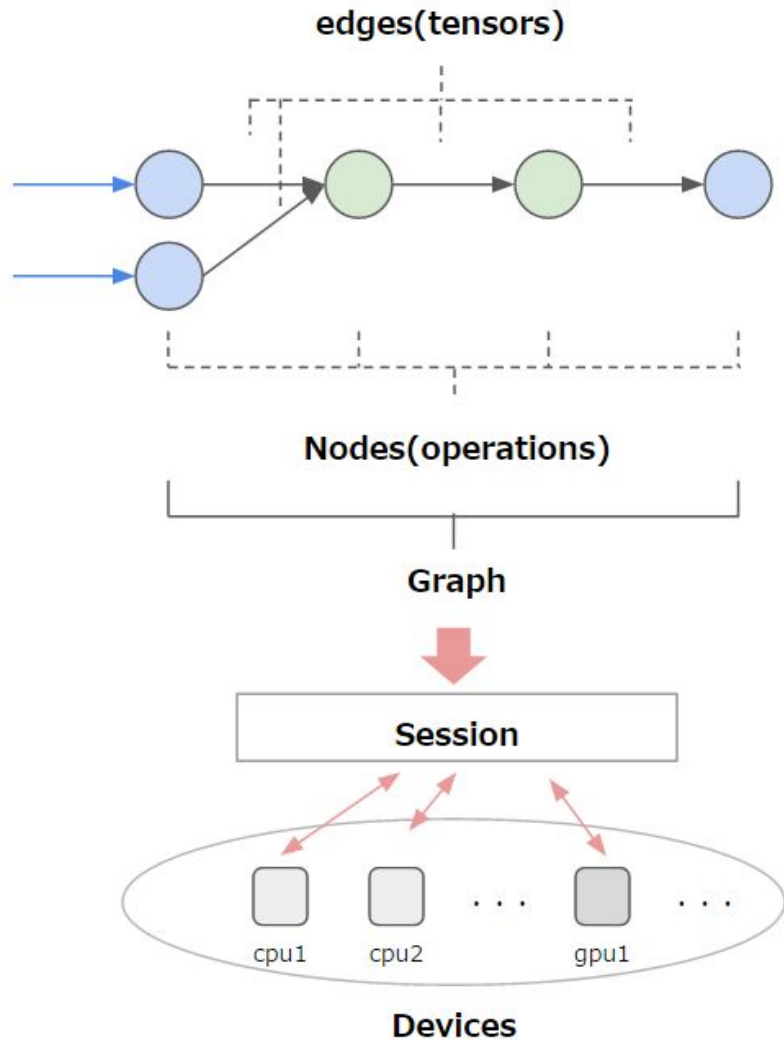
# Structuring TF Code 101



# Structuring TF Code 101



# Structuring TF Code 101



# Getting Stuff done in Tensorflow

Break it down into five steps

- Take in Inputs
- Define Network / Model
- Define Cost
- Define Optimization technique
- Train

**Try with MNIST example**

```
import tensorflow as tf
from tensorflow.examples.tutorials.mnist import
input_data
mnist = input_data.read_data_sets('MNIST_data',
                                  one_hot=True)
```

Placeholders

```
X = tf.placeholder(tf.float32, [128, 784])
Y_true = tf.placeholder(tf.float32, [128, 10])
```



```
import tensorflow as tf
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
```

```
X = tf.placeholder(tf.float32, [128, 784])
Y_true = tf.placeholder(tf.float32, [128, 10])
```

Parameters  
and  
Operations

```
m = tf.get_variable('m', [784, 10])
b = tf.get_variable('b', [10])
Y_pred = tf.nn.xw_plus_b(X, m, b)
```

```
X = tf.placeholder(tf.float32, [128, 784])
```

```
Y_true = tf.placeholder(tf.float32, [128, 10])
```

```
m = tf.get_variable('m', [784, 10])
```

```
b = tf.get_variable('b', [10])
```

```
Y_pred = tf.nn.xw_plus_b(X, m, b)
```

Cost

**cost =**

```
tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    Y_pred, Y_true))
```

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)

X = tf.placeholder(tf.float32, [128, 784])
Y_true = tf.placeholder(tf.float32, [128, 10])
m = tf.get_variable('m', [784, 10])
b = tf.get_variable('b', [10])
Y_pred = tf.nn.xw_plus_b(X, m, b)

cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(Y_pred,
                                                                Y_true))
```

Optimizer      **optimizer =**  
**tf.train.GradientDescentOptimizer(learning\_rate=0.5)**  
**.minimize(cost)**

```
X = tf.placeholder(tf.float32, [128, 784])
Y_true = tf.placeholder(tf.float32, [128, 10])
m = tf.get_variable('m', [784, 10])
b = tf.get_variable('b', [10])
Y_pred = tf.nn.xw_plus_b(X, m, b)

cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(Y_pred,
                                                                Y_true))

optimizer =
tf.train.GradientDescentOptimizer(learning_rate=0.5).minimize(cost)
```

## Train Code

```
sess = tf.Session()
sess.run(tf.initialize_all_variables())

for i in range(2000):
    trX, trY = mnist.train.next_batch(128)
    sess.run(optimizer, feed_dict={X: trX, Y_true: trY})
```




# Final Tricks

# Scaling to Optimize?


- When doing cost predictions using TF(Ex: Using softmax, similar to previous slide), pass in unscaled predictions, don't scale them first.
- TF does it's own optimization, and it does it brilliantly. If any scaling is to be done in order to avoid overflows/underflows, TF will take care of it

# Parameter Sharing

- In dropout and other train-time stochastic methods, make sure you use `tf.get_variable()` as giving different dropout keep probabilities actually creates two different models with two different variable sets. For instance, (see next few slides)



(a) Standard Neural Net




(b) After applying dropout.




# Without Weight Sharing

$Y\_pred = model(X, 0.8)$




$Y\_pred = model(X, 1.)$




# With Weight Sharing

$Y\_pred = model(X, 0.8)$



$Y\_pred = model(X, 1.)$



# Parameter Sharing

```
def model(X, p_keep):  
    m0 = tf.get_variable('m0', [784, 256])  
    b0 = tf.get_variable('b0', [256], initializer=tf.constant_initializer(0.))  
    m1 = tf.get_variable('m1', [256, 256])  
    b1 = tf.get_variable('b1', [256], initializer=tf.constant_initializer(0.))  
    m2 = tf.get_variable('m2', [256, 10])  
    b2 = tf.get_variable('b2', [10], initializer=tf.constant_initializer(0.))  
    h1 = tf.nn.relu(tf.nn.xw_plus_b(X, m0, b0))  
    h1 = tf.nn.dropout(h1, p_keep)  
    h2 = tf.nn.relu(tf.nn.xw_plus_b(h1, m1, b1))  
    h2 = tf.nn.dropout(h2, p_keep)  
    output = tf.nn.xw_plus_b(h2, m2, b2)  
    return output
```

```
Y_pred = model(X, 0.8)
```

```
Y_pred_test = model(X, 1.)
```

# Parameter Sharing (correct)

```
def model(X, p_keep):  
    m0 = tf.get_variable('m0', [784, 256])  
    b0 = tf.get_variable('b0', [256], initializer=tf.constant_initializer(0.))  
    m1 = tf.get_variable('m1', [256, 256])  
    b1 = tf.get_variable('b1', [256], initializer=tf.constant_initializer(0.))  
    m2 = tf.get_variable('m2', [256, 10])  
    b2 = tf.get_variable('b2', [10], initializer=tf.constant_initializer(0.))  
    h1 = tf.nn.relu(tf.nn.xw_plus_b(X, m0, b0))  
    h1 = tf.nn.dropout(h1, p_keep)  
    h2 = tf.nn.relu(tf.nn.xw_plus_b(h1, m1, b1))  
    h2 = tf.nn.dropout(h2, p_keep)  
    output = tf.nn.xw_plus_b(h2, m2, b2)  
    return output
```

**with tf.variable\_scope("model") as scope:**

**Y\_pred = model(X, 0.8)**

**scope.reuse\_variables()**

**Y\_pred\_test = model(X, 1.)**

# Collections

- Are a bucket which can contain parameters, operations, model predictions, etc for later reference more conveniently

P.S.: TF automatically creates some collections

# Collections

```
def model(X):  
    m0 = tf.get_variable('m0', [784, 256])  
    b0 = tf.get_variable('b0', [256], initializer=tf.constant_initializer(0.))  
    m1 = tf.get_variable('m1', [256, 256])  
    b1 = tf.get_variable('b1', [256], initializer=tf.constant_initializer(0.))  
    m2 = tf.get_variable('m2', [256, 10])  
    b2 = tf.get_variable('b2', [10], initializer=tf.constant_initializer(0.))  
    h1 = tf.nn.relu(tf.nn.xw_plus_b(X, m0, b0))  
    h2 = tf.nn.relu(tf.nn.xw_plus_b(h1, m1, b1))  
    tf.add_to_collection("activations", h1)  
    tf.add_to_collection("activations", h2)  
    output = tf.nn.xw_plus_b(h2, m2, b2)  
    return output
```

```
Y_pred = model(X)
```

# Collections

```
activations = tf.get_collection('activations')  
activations_values = session.run(activations)
```

```
parameters = tf.get_collection('trainable_parameters')  
parameter_values = session.run(parameters)
```

# Stuff to read up

- Image handling using TensorFlow:  
[https://www.tensorflow.org/versions/master/api\\_docs/python/image.html#encoding-and-decoding](https://www.tensorflow.org/versions/master/api_docs/python/image.html#encoding-and-decoding)
- How to run interactive sessions:
- TF Learn, a semi-high level API to TF
- TF Records, TF's default binary storage format



# Acknowledgements

- CS230i and CS231n - Stanford Courses
- NLintz' tutorials